

KNN Classification and Regression

#	Height (inches)	Weight (kgs)	B.P. Sys	B.P. Dia	Heart disease	Cholesterol Level
1	62	70	120	80	No	150
2	72	90	110	70	No	160
3	74	80	130	70	No	130
4	65	120	150	90	Yes	200
5	67	100	140	85	Yes	190
6	64	110	130	90	No	130
7	69	150	170	100	Yes	250
8	66	115	145	90		

KNN Classification and Regression

#	Height (inches)	Weight (kgs)	B.P. Sys	B.P. Dia	Heart disease	Cholesterol Level	Euclidean Distance
1	62	70	120	80	No	150	52.59
2	72	90	110	70	No	160	47.81
3	74	80	130	70	No	130	43.75
4	65	120	150	90	Yes	200	7.14
5	67	100	140	85	Yes	190	16.61
6	64	110	130	90	No	130	15.94
7	69	150	170	100	Yes	250	44.26
8	66	115	145	90			



Here the data in red is p and the other data is q, so the euclidean distance will be calculated as:
 $(p_1 - q_1)^2 + (p_2 - q_2)^2 + \dots$

So here::
 $(66 - 62)^2 + (115 - 70)^2 + (145 - 120)^2 + (90 - 80)^2$ ^1/2

Then the distance for s#1 will be calculated and same for others

Example: Handwritten digit recognition

- 16x16 bitmaps
- 8-bit grayscale
- Euclidean distances over raw pixels



- $$D(x, y) = \sqrt{\sum_{i=0}^{255} (x_i - y_i)^2}$$

Accuracy:

- 7-NN ~ 95.2%
- SVM ~ 95.8%
- Humans ~ 97.5%

The KNN Algorithm

Input: Training samples $D = \{(\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_n, y_n)\}$, Test sample $d = (\vec{x}, y)$, k . Assume \vec{x} to be an m -dimensional vector.

Output: Class label of test sample d

Algorithm/steps of kNN, most important

1. Compute the distance between d and every sample in D
2. Choose the K samples in D that are nearest to d ; denote the set by $S_d \in D$
3. Assign d the label y_i of the majority class in S_d

Note:

All action takes place in the test phase, the training phase is essentially to clean, normalize and store the data

Complexity of KNN

Input: Training samples $D = \{(\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_n, y_n)\}$, Test sample $d = (\vec{x}, y)$, k . Assume \vec{x} to be an m -dimensional vector.

Output: Class label of test sample d

1. Compute the distance between d and every sample in D

n samples, each is m -dimensional $\Rightarrow O(mn)$

2. Choose the K samples in D that are nearest to d ; denote the set by $S_d \in D$

- Either naively do K passes of all samples costing $O(n)$ each time for $O(nk)$
- Or use the *quickselect* algorithm (median of medians) to find the k th smallest distance in $O(n)$ and then return all distances no larger than the k th smallest distance. This will accumulate to $O(n)$

3. Assign d the label y_i of the majority class in S_d

This is $O(k)$.

Time complexity: $O(mn + n + k) = O(mn)$, assuming k to be a constant.

Choosing the value of K – The theory

k=1:

- High variance
- Small changes in the dataset will lead to big changes in classification
- Overfitting
- Is too specific and not well-generalized
- It tends to be sensitive to noise
- The model accomplishes a high accuracy on train set but will be a poor predictor on new, previously unseen data points

k= very large (e.g. 100):

- The model is too generalized and not a good predictor on both train and test sets.
- High bias
- Underfitting

k=n:

- The majority class in the dataset wins for every prediction
- High bias

Tuning the hyperparameter K – the Method

- Divide your training data into training and validation sets.
- Do multiple iterations of m-fold cross-validation, each time with a different value of k, starting from k=1
- Keep iterating until the k with the best classification accuracy (minimal loss) is found
- What happens if we use the training set itself, instead of a validation set? Which k wins?
 - K=1, as there is always a nearest instance with the correct label, the instance itself

KNN – The good, the bad and the ugly

KNN is a simple algorithm but is highly effective for solving various real life classification problems. Especially when the datasets are large and continuously growing.

We will show that as $n \rightarrow \infty$, the 1-NN classifier is only a factor 2 worse than the best possible classifier (remember our old friend the Bayes Optimal Classifier?).

Challenges:

1. How to find the optimum value of K?
2. How to find the right distance function?

Problems:

1. High computational time cost for each prediction.
2. High memory requirement as we need to keep all training samples.
3. The curse of dimensionality.

Bayes Error https://www.cs.cornell.edu/courses/cs4780/2018fa/lectures/lecturenote02_kNN.html

Assume that we know $P(y|x)$, so we can simply predict the correct label y^* as:

$$y^* = h_{opt}(x) = \operatorname{argmax}_y P(y|x)$$

Although the Bayes Classifier is optimal, but it is not perfect and can still make mistakes. E.g. It will predict incorrectly when a test point does not have the most likely label.

So, if $P(y^*|x)$ is the probability of correct classification then the probability of incorrect classification, the Bayes Error is given as:


$$\epsilon_{Bayes} = 1 - P(y^*|x)$$

1-NN Error as $n \rightarrow \infty$ (Cover and Hart 1967, Weinberger Lec 2)


Let x_{NN} be the nearest neighbor of our test point x_t

As $n \rightarrow \infty$, $\text{dist}(x_{NN}, x_t) \rightarrow 0$


- i.e. $x_{NN} \rightarrow x_t$
- 1-NN returns the label of x_{NN}



Small n



Large n



$n \rightarrow \infty$

What is the probability that this is not the correct label of x_t ?

- As $x_{NN} \rightarrow x_t$, the probability of misclassification is the same as the probability of x_{NN} and x_t having different labels

1-NN Error as $n \rightarrow \infty$ (Cover and Hart 1967, Weinberger Lec 2)

There are two ways this could happen.

- What's the probability that x_{NN} had the correct label y^* ?
 - $P(y^*|x_{NN})$
- What's the probability that x_{NN} did not have the correct label y^* ?
 - $1 - P(y^*|x_{NN})$
- What's the probability that y^* was the correct label of x_t ?
 - $P(y^*|x_t)$

1. So, what's the probability that y^* was the correct label of x_t but the nearest neighbor x_{NN} did not have that label?

$$P(y^*|x_t)(1 - P(y^*|x_{NN}))$$

- What's the probability that y^* was the correct label of x_t ?
 - $P(y^*|x_t)$
- What's the probability that y^* was not the correct label of x_t ?
 - $1 - P(y^*|x_t)$
- What's the probability that x_{NN} had the correct label y^* ?
 - $P(y^*|x_{NN})$

2. So, what's the probability that y^* was not the correct label of x_t but the nearest neighbor x_{NN} had that label?

$$P(y^*|x_{NN})(1 - P(y^*|x_t))$$

1-NN Error as $n \rightarrow \infty$ (Cover and Hart 1967, Weinberger Lec 2)

So, the total probability of misclassification is:

$$\epsilon_{NN} = P(y^*|x_t)(1 - P(y^*|x_{NN})) + P(y^*|x_{NN})(1 - P(y^*|x_t))$$

As $P(y^*|x_t) \leq 1$ and $P(y^*|x_{NN}) \leq 1$,

$$\begin{aligned}\epsilon_{NN} &= P(y^*|x_t)(1 - P(y^*|x_{NN})) + P(y^*|x_{NN})(1 - P(y^*|x_t)) \\ &\leq 1(1 - P(y^*|x_{NN})) + 1(1 - P(y^*|x_t))\end{aligned}$$

As, $x_{NN} \rightarrow x_t$, $P(y^*|x_{NN}) = P(y^*|x_t)$

$$\epsilon_{NN} \leq (1 - P(y^*|x_t)) + (1 - P(y^*|x_t))$$

$$\epsilon_{NN} \leq 2(1 - P(y^*|x_t))$$

$$\epsilon_{NN} \leq 2\epsilon_{Bayes}$$